

▼ Implementing Feed-forward neural networks with Keras and TensorFlow

a. Import the necessary packages

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
```

b. Load the training and testing data (MNIST/CIFAR10)

```
# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

#normalize the input data to range [0, 1]
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

#flatten the input data (28 X 28 images to 784 pixels)
x_train = x_train.reshape(-1, 784)
x_test = x_test.reshape(-1, 784)

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
```

c. Define the network architecture using Keras

```
model = Sequential()
model.add(Dense(128, activation='relu', input_shape=(784,)))
model.add(Dense(64, activation='relu'))
model.add(Dense(10, activation='softmax'))
```

d. Train the model using SGD

```
# from keras.src.engine.training import optimizer
# compile the model
model.compile(optimizer='sgd', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=64, validation_split=0.2)
```

```
Epoch 1/10
750/750 [=====] - 7s 3ms/step - loss: 1.0122 - accuracy: 0.7379 - val_loss: 0.4530 - val_accuracy: 0.8848
Epoch 2/10
750/750 [=====] - 3s 4ms/step - loss: 0.4081 - accuracy: 0.8876 - val_loss: 0.3333 - val_accuracy: 0.9087
Epoch 3/10
750/750 [=====] - 2s 3ms/step - loss: 0.3307 - accuracy: 0.9064 - val_loss: 0.2933 - val_accuracy: 0.9152
Epoch 4/10
750/750 [=====] - 3s 3ms/step - loss: 0.2942 - accuracy: 0.9165 - val_loss: 0.2718 - val_accuracy: 0.9220
Epoch 5/10
750/750 [=====] - 2s 3ms/step - loss: 0.2690 - accuracy: 0.9232 - val_loss: 0.2534 - val_accuracy: 0.9277
Epoch 6/10
750/750 [=====] - 2s 3ms/step - loss: 0.2501 - accuracy: 0.9285 - val_loss: 0.2336 - val_accuracy: 0.9338
Epoch 7/10
750/750 [=====] - 3s 4ms/step - loss: 0.2339 - accuracy: 0.9332 - val_loss: 0.2220 - val_accuracy: 0.9369
Epoch 8/10
750/750 [=====] - 3s 4ms/step - loss: 0.2198 - accuracy: 0.9367 - val_loss: 0.2105 - val_accuracy: 0.9407
Epoch 9/10
750/750 [=====] - 3s 3ms/step - loss: 0.2073 - accuracy: 0.9408 - val_loss: 0.2015 - val_accuracy: 0.9447
Epoch 10/10
750/750 [=====] - 2s 3ms/step - loss: 0.1963 - accuracy: 0.9441 - val_loss: 0.1923 - val_accuracy: 0.9463
```

e. Evaluate the network

```
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_accuracy:.4f}')

313/313 [=====] - 1s 2ms/step - loss: 0.1941 - accuracy: 0.9451
Test Loss: 0.1941
Test Accuracy: 0.9451
```



f. Plot the training loss and accuracy

```
train_loss = history.history['loss']
val_loss = history.history['val_loss']
train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(1, len(train_loss) + 1)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(epochs, train_loss, 'b', label='Training Loss')
plt.plot(epochs, val_loss, 'r', label='Validation Loss')
plt.title('Training and validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(epochs, train_acc, 'b', label='Training Accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```



▼ Build the Image classification model by dividing the model into the following four stages:

a. Loading and preprocessing the image data

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.utils import to_categorical

# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Preprocess the image data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

#one hot encoding
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 3s 0us/step
```

b. Defining the model's architecture

```
# Define the model architecture
model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(10, activation='softmax'))
```

c. Training the model

```
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=64, validation_data=(x_test, y_test))

Epoch 1/10
782/782 [=====] - 12s 6ms/step - loss: 1.5589 - accuracy: 0.4357 - val_loss: 1.2763 - val_accuracy: 0.5453
Epoch 2/10
782/782 [=====] - 4s 5ms/step - loss: 1.1791 - accuracy: 0.5815 - val_loss: 1.0912 - val_accuracy: 0.6115
Epoch 3/10
782/782 [=====] - 5s 7ms/step - loss: 1.0284 - accuracy: 0.6385 - val_loss: 1.0530 - val_accuracy: 0.6310
Epoch 4/10
782/782 [=====] - 4s 5ms/step - loss: 0.9383 - accuracy: 0.6691 - val_loss: 0.9637 - val_accuracy: 0.6710
Epoch 5/10
782/782 [=====] - 5s 7ms/step - loss: 0.8625 - accuracy: 0.6981 - val_loss: 0.9030 - val_accuracy: 0.6923
Epoch 6/10
782/782 [=====] - 4s 6ms/step - loss: 0.8039 - accuracy: 0.7182 - val_loss: 0.9372 - val_accuracy: 0.6791
Epoch 7/10
782/782 [=====] - 4s 6ms/step - loss: 0.7536 - accuracy: 0.7349 - val_loss: 0.9106 - val_accuracy: 0.6888
Epoch 8/10
782/782 [=====] - 5s 6ms/step - loss: 0.7048 - accuracy: 0.7536 - val_loss: 0.8679 - val_accuracy: 0.7083
Epoch 9/10
782/782 [=====] - 4s 6ms/step - loss: 0.6709 - accuracy: 0.7664 - val_loss: 0.8758 - val_accuracy: 0.7097
Epoch 10/10
782/782 [=====] - 4s 6ms/step - loss: 0.6265 - accuracy: 0.7795 - val_loss: 0.8569 - val_accuracy: 0.7179
```

d. Estimating the model's performance

```
# Evaluate the model
test_loss, test_accuracy = model.evaluate(x_test, y_test)

print(f'Test loss: {test_loss:.4f}')
print(f'Test accuracy: {test_accuracy:.4f}')

313/313 [=====] - 1s 3ms/step - loss: 0.8569 - accuracy: 0.7179
Test loss: 0.8569
Test accuracy: 0.7179
```

▼ Use Autoencoder to implement anomaly detection. Build the model by using the following:

a. Import required libraries

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Input
from tensorflow.keras.datasets import mnist
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

b. Upload/access the dataset

```
# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Preprocess the image data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

#split data into training and testing
x_train, x_val = train_test_split(x_train, test_size=0.2, random_state=42)
```

c. The encoder converts it into a latent representation

```
input_layer = Input(shape=(x_train.shape[1],))
encoded = Dense(128, activation='relu')(input_layer)
encoded = Dense(64, activation='relu')(encoded)
encoded = Dense(32, activation='relu')(encoded)

encoder = Model(input_layer, encoded)
```

d. Decoder networks convert it back to the original input

```
decoded = Dense(64, activation='relu')(encoded)
decoded = Dense(128, activation='relu')(decoded)
decoded = Dense(x_train.shape[1], activation='sigmoid')(decoded)

autoencoder = Model(input_layer, decoded)
```

e. Compile the models with Optimizer, Loss, and Evaluation Metrics

```
# Compile the encoder model
autoencoder.compile(optimizer='adam', loss='mean_squared_error')

#train the autoencoder on the training data
autoencoder.fit(x_train, x_train, epochs=10, batch_size=32, validation_data=(x_val, x_val))
```

```
Epoch 1/10
1500/1500 [=====] - 8s 4ms/step - loss: 0.0359 - val_loss: 0.0215
Epoch 2/10
1500/1500 [=====] - 6s 4ms/step - loss: 0.0188 - val_loss: 0.0169
Epoch 3/10
1500/1500 [=====] - 5s 4ms/step - loss: 0.0155 - val_loss: 0.0146
Epoch 4/10
1500/1500 [=====] - 6s 4ms/step - loss: 0.0138 - val_loss: 0.0134
Epoch 5/10
1500/1500 [=====] - 5s 4ms/step - loss: 0.0127 - val_loss: 0.0125
```

```

Epoch 6/10
1500/1500 [=====] - 6s 4ms/step - loss: 0.0121 - val_loss: 0.0120
Epoch 7/10
1500/1500 [=====] - 5s 4ms/step - loss: 0.0115 - val_loss: 0.0115
Epoch 8/10
1500/1500 [=====] - 6s 4ms/step - loss: 0.0111 - val_loss: 0.0111
Epoch 9/10
1500/1500 [=====] - 6s 4ms/step - loss: 0.0107 - val_loss: 0.0107
Epoch 10/10
1500/1500 [=====] - 6s 4ms/step - loss: 0.0104 - val_loss: 0.0107
<keras.src.callbacks.History at 0x799e739f4b80>

```

```

#evaluate the autoencoder on the training data
reconstruction_error = np.mean(np.square(x_test - autoencoder.predict(x_test)), axis=1)

```

```

#set the threshold anomaly detection
threshold = np.percentile(reconstruction_error, 95)

```

```

#detect the anomaly in the test data
anomalies = x_test[reconstruction_error > threshold]

```

```


313/313 [=====] - 1s 2ms/step

```

```

#display the some detected anomalies
n_anomalies_to_display = 5
for i in range(n_anomalies_to_display):
    plt.subplot(1, n_anomalies_to_display, i + 1)
    plt.imshow(anomalies[i].reshape(28, 28), cmap='gray')
    plt.title(f'Anomaly {i + 1}')
    plt.axis('off')
plt.show()

```



▼ Implement the Continuous Bag of Words (CBOW) Model. Stages can be:

a. Data preparation

```
corpus = [  
    "I like to learn deep learning",  
    "Deep learning is interesting",  
    "I enjoy studying deep learning"  
]
```

b. Generate training data

```
from tensorflow.keras.preprocessing.text import Tokenizer  
from tensorflow.keras.preprocessing.sequence import skipgrams  
  
# Tokenize the corpus  
tokenizer = Tokenizer()  
tokenizer.fit_on_texts(corpus)  
word2idx = tokenizer.word_index  
idx2word = {v: k for k, v in word2idx.items()}  
  
# Generate training data  
vocab_size = len(word2idx) + 1  
target_words, context_words = [], []  
for sentence in corpus:  
    tokenized = tokenizer.texts_to_sequences([sentence])[0]  
    pairs, _ = skipgrams(tokenized, vocabulary_size=vocab_size, window_size=1, negative_samples=0)  
    target, context = zip(*pairs)  
    target_words.extend(target)  
    context_words.extend(context)
```

c. Train model

```
import tensorflow as tf  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Embedding, Dense, Reshape  
  
embedding_dim = 100  
  
model = Sequential()  
model.add(Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=1))  
model.add(Reshape((embedding_dim,)))  
model.add(Dense(units=vocab_size, activation='softmax'))  
  
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam')  
model.summary()  
  
# Train the model  
model.fit(x=target_words, y=context_words, epochs=50)
```

```

1/1 [=====] - 0s 10ms/step - loss: 2.1081
Epoch 31/50
1/1 [=====] - 0s 11ms/step - loss: 2.0972
Epoch 32/50
1/1 [=====] - 0s 12ms/step - loss: 2.0861
Epoch 33/50
1/1 [=====] - 0s 14ms/step - loss: 2.0750
Epoch 34/50
1/1 [=====] - 0s 11ms/step - loss: 2.0638
Epoch 35/50
1/1 [=====] - 0s 11ms/step - loss: 2.0524
Epoch 36/50
1/1 [=====] - 0s 12ms/step - loss: 2.0410
Epoch 37/50
1/1 [=====] - 0s 11ms/step - loss: 2.0294
Epoch 38/50
1/1 [=====] - 0s 12ms/step - loss: 2.0178
Epoch 39/50
1/1 [=====] - 0s 12ms/step - loss: 2.0060
Epoch 40/50
1/1 [=====] - 0s 11ms/step - loss: 1.9941
Epoch 41/50
1/1 [=====] - 0s 13ms/step - loss: 1.9821
Epoch 42/50
1/1 [=====] - 0s 12ms/step - loss: 1.9700
Epoch 43/50
1/1 [=====] - 0s 13ms/step - loss: 1.9578
Epoch 44/50
1/1 [=====] - 0s 13ms/step - loss: 1.9455
Epoch 45/50
1/1 [=====] - 0s 11ms/step - loss: 1.9331
Epoch 46/50
1/1 [=====] - 0s 11ms/step - loss: 1.9206
Epoch 47/50
1/1 [=====] - 0s 11ms/step - loss: 1.9080
Epoch 48/50
1/1 [=====] - 0s 19ms/step - loss: 1.8953
Epoch 49/50
1/1 [=====] - 0s 12ms/step - loss: 1.8825
Epoch 50/50
1/1 [=====] - 0s 11ms/step - loss: 1.8696
<keras.src.callbacks.History at 0x799e7d7e3d30>

```

d. Output

```

word_to_lookup = "deep"
word_idx = word2idx[word_to_lookup]
word_embedding = model.layers[0].get_weights()[0][word_idx]

print(f"Embedding for '{word_to_lookup}': {word_embedding}")

```

```

Embedding for 'deep': [-0.08426171 -0.04487503  0.03471786  0.00877153  0.03311342 -0.05220547
 -0.10537343  0.05674294  0.064823   -0.0503858  0.08464649 -0.03675745
  0.09589159 -0.02757626  0.05152629  0.02986002  0.00835305  0.0891538
  0.0274698  -0.08283201  0.08814979 -0.02207836  0.07251184 -0.10077845
  0.01189684 -0.04852607 -0.03135568 -0.0921233  0.0509307  0.0990841
 -0.07088152 -0.05199756 -0.08785651  0.02239579  0.07870483 -0.08659703
  0.06364825 -0.0147792  -0.07328346 -0.00406267 -0.06556738 -0.05483082
  0.02189518 -0.06745288  0.06877382 -0.08121266 -0.00131394 -0.08789965
  0.01931682 -0.06161145  0.04886431  0.01397459  0.01659805  0.0495768
 -0.00831992 -0.06807461  0.0378589  -0.10111599  0.06906442  0.07710025
  0.01673487  0.05439958 -0.02330448  0.05384533 -0.05315861  0.09971622
 -0.02362644  0.02694047  0.00438797  0.00650261 -0.04028632  0.08963581
 -0.08902054 -0.00401343 -0.05719236  0.07904532  0.09287891  0.06366136
  0.07346602  0.00621467  0.04426254 -0.01764723 -0.00551786 -0.05514027
 -0.04602342  0.05532861 -0.07229256 -0.07894777  0.03305203  0.04546653
  0.07533833  0.05800479 -0.05717856  0.02284533 -0.09087887  0.07040332
 -0.03173155 -0.07509761 -0.0395535  0.05414165]

```


▼ Object detection using Transfer Learning of CNN architectures

```
import torch
import torchvision
from torchvision import transforms as T
from PIL import Image
import cv2
from google.colab.patches import cv2_imshow

# Load the pre-trained model
model = torchvision.models.detection.fasterrcnn_resnet50_fpn(pretrained=True)

/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None`
warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth" to /root/.cache/torch/hub/checkpoints/
100%|██████████| 160M/160M [00:01<00:00, 160MB/s]

model.eval()

!wget 'http://images.cocodataset.org/val2017/000000037777.jpg'

--2023-10-19 13:34:44-- http://images.cocodataset.org/val2017/000000037777.jpg
Resolving images.cocodataset.org (images.cocodataset.org)... 52.216.49.97, 3.5.25.75, 3.5.7.189, ...
Connecting to images.cocodataset.org (images.cocodataset.org)|52.216.49.97|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 40833 (40K) [image/jpeg]
Saving to: '000000037777.jpg'

000000037777.jpg 100%[=====] 39.88K --.-KB/s in 0.05s

2023-10-19 13:34:45 (842 KB/s) - '000000037777.jpg' saved [40833/40833]

ig = Image.open('/content/000000037777.jpg')

transform = T.ToTensor()
img = transform(ig)

with torch.no_grad():
    pred = model([img])

pred[0].keys()

dict_keys(['boxes', 'labels', 'scores'])

bboxex, labels, scores = pred[0]['boxes'], pred[0]['labels'], pred[0]['scores']

num = torch.argmax(scores > 0.8).shape[0]

print(num)

11

coco_names = [
    "person", "bicycle", "car", "motorcycle", "airplane", "bus", "train", "truck", "boat", "traffic light",
    "fire hydrant", "street sign", "stop sign", "parking meter", "bench", "bird", "cat", "dog", "horse", "sheep",
    "cow", "elephant", "bear", "zebra", "giraffe", "hat", "backpack", "umbrella", "shoe", "eye glasses", "handbag",
    "tie", "suitcase", "frisbee", "skis", "snowboard", "sports ball", "kite", "baseball bat",
    "baseball glove", "skateboard", "surfboard", "tennis racket", "bottle",
    "plate", "wine glass", "cup", "fork", "knife", "spoon", "bowl",
    "banana", "apple", "sandwich", "orange", "broccoli", "carrot", "hot dog",
    "pizza", "donut", "cake", "chair", "couch", "potted plant", "bed",
    "mirror", "dining table", "window", "desk", "toilet", "door", "tv",
    "laptop", "mouse", "remote", "keyboard", "cell phone", "microwave",
    "oven", "toaster", "sink", "refrigerator", "blender", "book",
```

```

"clock" , "vase" , "scissors" , "teddy bear" , "hair drier" , "toothbrush" , "hair brush"
]

font = cv2.FONT_HERSHEY_SIMPLEX

igg = cv2.imread('/content/000000037777.jpg')
for i in range(num):
    x1, y1, x2, y2 = bboxex[i].numpy().astype('int')
    class_name = coco_names[labels.numpy()[i] - 1]
    igg = cv2.rectangle(igg, (x1, y1), (x2, y2), (255, 0, 0), 1)
    igg = cv2.putText(igg, class_name, (x1, y1-10), font, 0.5, (255, 0, 0), 1, cv2.LINE_AA)

cv2.imshow(igg)

```

